

# C(++) or Python? Cython!


Make code run up to 1000 x faster in only 5 minutes

Arne F. Meyer

Gatsby/SWC PyClub

16th January 2018

# Background



- C(++) is very fast but often inconvenient for research (especially plotting)
- Interpreted languages (here: Python) are excellent for research but in some cases very slow
- Common procedure: Where speed is needed, use a compiled language, then wrap the code for use from Python

### Donald Knuth (1974)


“Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We should forget about small efficiencies, say about 97% of the time: **premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%.**”

### Donald Knuth (1974)

“Programmers waste enormous amounts of time thinking about, or worrying about, the speed of noncritical parts of their programs, and these attempts at efficiency actually have a strong negative impact when debugging and maintenance are considered. We should forget about small efficiencies, say about 97% of the time: **premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%.**”


- Re-writing code is often very time-consuming and prone to errors
- How to optimize the critical 3% efficiently?

# Background



- C(++) is very fast but often inconvenient for research (especially plotting)
- Interpreted languages (here: Python) are excellent for research but in some cases very slow
- Common procedure: Where speed is needed, use a compiled language, then wrap the code for use from Python

# Background



- C(++) is very fast but often inconvenient for research (especially plotting)
- Interpreted languages (here: Python) are excellent for research but in some cases very slow
- Common procedure: Where speed is needed, use a compiled language, then wrap the code for use from Python
- Cython: combines the best of both worlds

# Cython at a glance

- Open-source project: [www.cython.org](http://www.cython.org)
- An optimizing compiler for the Python language
- Very active development
- Rapidly growing user base (many from science)

## Use-cases:

- ① Compiling Python code to machine-code
  - Supports a big subset of the Python language
  - Runs about 30% faster than plain Python code
- ② Add types for speedups (hundreds of times)
  - Optimize, don't re-write!
- ③ Easily use native libraries (C/C++/Fortran) directly
  - There are better tools, e.g., SWIG

# Example

## Ridge regression using stochastic gradient descent

Goal: minimize

$$\frac{1}{2} \sum_i (y_i - \mathbf{x}_i^T \mathbf{w})^2 + \frac{1}{2} \alpha \|\mathbf{w}\|^2$$

Pseudo code:

**input:**  $\{\mathbf{x}_i, y_i\}$ ,  $\alpha$ ,  $N_{\text{iter}}$

$\mathbf{w} \leftarrow \mathbf{0}$

**for**  $t = 1, 2, \dots, N_{\text{iter}}$  **do**

$\mathbf{x}_i, y_i \leftarrow$  draw random sample

$\gamma \leftarrow \frac{1}{\alpha t}$

$\mathbf{w} \leftarrow \mathbf{w} - \gamma \alpha \mathbf{w}$

$\mathbf{w} \leftarrow \mathbf{w} - \gamma \mathbf{x}_i^T (y_i - \mathbf{x}_i^T \mathbf{w})$

**end**



# Naive Python implementation

```
def ridge_sgd_naive(X, y, w, alpha, perm):  
  
    D = X.shape[1]  
    for t, i in enumerate(perm):  
  
        gamma = 1. / (1 + alpha*t)  
  
        # regularization step  
        for j in range(D):  
            w[j] *= (1. - gamma * alpha)  
  
        # loss step  
        z = 0  
        for j in range(D):  
            z += w[j] * X[i, j]  
  
        for j in range(D):  
            w[j] += gamma * X[i, j] * (z - y[i])
```

# Naive Python implementation

```
def ridge_sgd_naive(X, y, w, alpha, perm):  
  
    D = X.shape[1]  
    for t, i in enumerate(perm):  
  
        gamma = 1. / (1 + alpha*t)  
  
        # regularization step  
        for j in range(D):  
            w[j] -= (1. - gamma * alpha)  
  
        # loss step  
        z = 0  
        for j in range(D):  
            z += w[j] * X[i, j]  
  
        for j in range(D):  
            w[j] += gamma * X[i, j] * (z - y[i])
```

- Python: approx. 135 s

- Cython: approx. 97 s

```
import pyximport  
pyximport.install()
```

```
from cython_file import cython_function  
...
```

# Vectorized (Numpy) implementation

```
import numpy as np

def ridge_sgd_vectorized(X, y, w, alpha, perm):

    for t, i in enumerate(perm):

        gamma = 1. / (1 + alpha*t)

        # regularization step
        w *= (1. - gamma * alpha)

        # loss step
        z = np.dot(w, X[i, :])
        w += gamma * X[i, :] * (z - y[i])
```

# Vectorized (Numpy) implementation

```
import numpy as np

def ridge_sgd_vectorized(X, y, w, alpha, perm):

    for t, i in enumerate(perm):

        gamma = 1. / (1 + alpha*t)

        # regularization step
        w *= (1. - gamma * alpha)

        # loss step
        z = np.dot(w, X[i, :])
        w += gamma * X[i, :] * (z - y[i])
```

- Python: approx. 1.65 s
- Cython: approx. 1.44 s

# Cython: adding static types to naive implementation

```
def ridge_sgd_cython_types(np.ndarray[np.float64_t, ndim=2] X,
                           np.ndarray[np.float64_t, ndim=1] y,
                           np.ndarray[np.float64_t, ndim=1] w, double alpha,
                           np.ndarray[np.int64_t, ndim=1] perm):

    cdef int D = X.shape[1]
    cdef int i, j, t
    cdef double gamma, z

    for t, i in enumerate(perm):

        gamma = 1. / (1. + alpha*t)

        # regularization step
        for j in range(D):
            w[j] -= (1. - gamma * alpha)

        # loss step
        z = 0
        for j in range(D):
            z += w[j] * X[i, j]

        for j in range(D):
            w[j] += gamma * X[i, j] * (z - y[i])
```

# Cython: adding static types to naive implementation

```
def ridge_sgd_cython_types(np.ndarray[np.float64_t, ndim=2] X,
                           np.ndarray[np.float64_t, ndim=1] y,
                           np.ndarray[np.float64_t, ndim=1] w, double alpha,
                           np.ndarray[np.int64_t, ndim=1] perm):

    cdef int D = X.shape[1]
    cdef int i, j, t
    cdef double gamma, z

    for t, i in enumerate(perm):

        gamma = 1. / (1. + alpha*t)

        # regularization step
        for j in range(D):
            w[j] -= (1. - gamma * alpha)

        # loss step
        z = 0
        for j in range(D):
            z += w[j] * X[i, j]

        for j in range(D):
            w[j] += gamma * X[i, j] * (z - y[i])
```

- Run time: approx. 0.33 s

# Cython: static types + memoryviews

```
def ridge_sgd_cython_types(double[:, ::1] X,  
                           double[:] y,  
                           double[:] w, double alpha,  
                           long[:] perm):  
  
    cdef int D = X.shape[1]  
    cdef int i, j, t  
    cdef double gamma, z  
  
    for t, i in enumerate(perm):  
  
        gamma = 1. / (1. + alpha*t)  
  
        # regularization step  
        for j in range(D):  
            w[j] *= (1. - gamma * alpha)  
  
        # loss step  
        z = 0  
        for j in range(D):  
            z += w[j] * X[i, j]  
  
        for j in range(D):  
            w[j] += gamma * X[i, j] * (z - y[i])
```

- <http://cython.readthedocs.io/en/latest/src/userguide/memoryviews.html>

# Cython: static types + memoryviews

```
def ridge_sgd_cython_types(double[:, ::1] X,  
                           double[:] y,  
                           double[:] w, double alpha,  
                           long[:] perm):  
  
    cdef int D = X.shape[1]  
    cdef int i, j, t  
    cdef double gamma, z  
  
    for t, i in enumerate(perm):  
  
        gamma = 1. / (1. + alpha*t)  
  
        # regularization step  
        for j in range(D):  
            w[j] *= (1. - gamma * alpha)  
  
        # loss step  
        z = 0  
        for j in range(D):  
            z += w[j] * X[i, j]  
  
        for j in range(D):  
            w[j] += gamma * X[i, j] * (z - y[i])
```

- <http://cython.readthedocs.io/en/latest/src/userguide/memoryviews.html>
- Run time: approx. 0.31 s



## Cython: static types and C pointers

```
def ridge_sgd_cython_pointers(np.ndarray[np.float64_t, ndim=2] X,  
                              np.ndarray[np.float64_t, ndim=1] y,  
                              np.ndarray[np.float64_t, ndim=1] w, double alpha,  
                              np.ndarray[np.int64_t, ndim=1] perm):  
  
    cdef int D = X.shape[1]  
    cdef int i, j, t  
    cdef double gamma, z  
  
    cdef double *Xp = <double*> X.data  
    cdef double *yp = <double*> y.data  
    cdef double *wp = <double*> w.data  
    cdef long *pp = <long*> perm.data  
  
    for t, i in enumerate(perm):  
  
        ...  
  
        for j in range(D):  
            z += wp[j] * Xp[i*D + j]  
  
        for j in range(D):  
            wp[j] += gamma * Xp[i*D + j] * (z - yp[i])
```

# Cython: static types and C pointers

```
def ridge_sgd_cython_pointers(np.ndarray[np.float64_t, ndim=2] X,
                              np.ndarray[np.float64_t, ndim=1] y,
                              np.ndarray[np.float64_t, ndim=1] w, double alpha,
                              np.ndarray[np.int64_t, ndim=1] perm):

    cdef int D = X.shape[1]
    cdef int i, j, t
    cdef double gamma, z

    cdef double *Xp = <double*> X.data
    cdef double *yp = <double*> y.data
    cdef double *wp = <double*> w.data
    cdef long *pp = <long*> perm.data

    for t, i in enumerate(perm):
        ...

        for j in range(D):
            z += wp[j] * Xp[i*D + j]

        for j in range(D):
            wp[j] += gamma * Xp[i*D + j] * (z - yp[i])
```

- Run time: approx. 0.24 s
- Replacing loops by BLAS functions: approx. 0.18 s

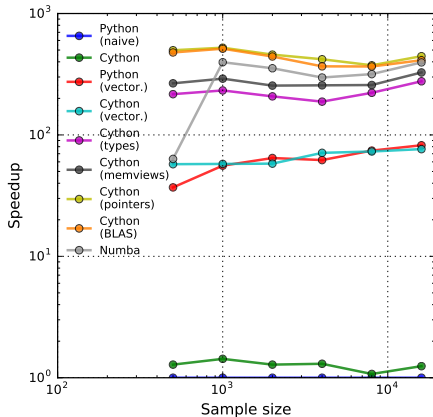

# Numba JIT implementation

```
from numba.decorators import autojit  
  
ridge_sgd_numba = autojit(ridge_sgd_vectorized)
```

- <http://numba.pydata.org/>
- Just-in-time (JIT) compiler
- Run time: approx. 0.22 s

# Summary


## Stochastic gradient descent



- Cython about 100 – 500 times faster than naive Python
- Cython about 5 – 10 times faster than (vectorized) Numpy
- Comparable to Numba

## Example 2



### Recurrent neural network



- Loebel & Tsodyks (2007)
- 15 coupled EI networks (cortical columns)
- Each column:  $N_E = 100$ ,  $N_I = 100$
- External stimulus input

## Example 2

### Recurrent neural network



- Cython about 85 times faster than naive Python
- Cython about 7 times faster than (vectorized) Numpy
- JIT compiler (Numba) much slower than Cython version

# What was that all about?



- Goal: writing fast code in interpreted language
- Avoid unnecessary re-writing of (working) code
- Cython: simply add static types to existing (Python) code
- Only a few extra lines (about 5 minutes ...)
- Speedup: 50-1000 times (naive Python), 1-250 times (vectorized Numpy)
- In some cases, JIT compilers (e.g., Numba) may help, too